

LEADER ELECTION IN RAFT

1 Implementation of the Election Algorithm

1.1 System Initialization and Heartbeats

In the Raft consensus algorithm, servers operate in one of the following states [1, 2]:

- Follower
- Candidate
- Leader

All servers are initialized as Followers [1, 2]. A Leader maintains its authority by broadcasting periodic heartbeat messages, through empty `AppendEntries` RPCs [1, 2].

1.2 Initiating an Election

If a Follower does not receive a heartbeat within a specific timeframe, known as the *"election timeout"*, it assumes the Leader has failed [1, 2]. At this point, the Follower transitions to the Candidate state, increments its current term number, votes for itself, and broadcasts `RequestVote` RPCs to the other servers in the cluster [1, 2].

1.3 Processing Vote Requests

Normally, a `RequestVote` RPC contains the candidate's `lastLogIndex` and `lastLogTerm` to allow voters to determine which candidate's log is more up-to-date [1]. However, the assignment description specifies that *"all `AppendEntries` RPC carry empty log messages and no node ever appends anything to its log"*. In this particular case, the evaluation of the `RequestVote` RPC will be based only on the term number, rejecting the request if the candidate's term is lower than its own, and giving the vote if it has not already voted in the current term [2].

2 Election Mechanics with Empty Logs

So in this case the leader election follows a first-come-first-served logic, based on term numbers. The Raft protocol dictates that a candidate may only be elected if its log is at least as up-to-date as any other log in that majority [1]. In our example, the logs are empty and thus, identical, so this condition is automatically respected.

2.1 Securing Leadership

Thus, a follower will grant its vote to the first candidate that requests it during a new term [2]. Upon securing votes from a majority of the servers, including its own vote, a candidate promotes itself to the leader state. It then immediately dispatches heartbeat messages to assert its leadership and suppress further elections [2]. If a candidate receives a heartbeat from an entity claiming to be leader with a term number equal to or greater than its own, it acknowledges the leader and reverts to the follower state [2].

3 Vulnerability to Election Failures

3.1 Split Votes

Despite the simplified state of the logs, elections can still fail due to split votes [2]. If multiple followers experience an election timeout simultaneously, they will transition into candidates, increment their term numbers, and cast votes for themselves at the exact same time [2]. Because the remaining servers allocate their votes on a first-come-first-served basis, the votes may be distributed in such a way that no single candidate achieves the requisite majority [2].

3.2 Mitigation with Randomized Timeouts

This results in a failed election and a stalemate that persists until a candidate timeout occurs [2]. At this point, the candidates increment their term numbers and initiate a subsequent election round [2]. To mitigate the risk of indefinite split-vote loops, Raft employs randomized election timeouts on a per-server basis [2]. By assigning timeouts randomly from a fixed interval, such as 150 to 300 milliseconds, the protocol ensures that it is highly probable for one server to time out before its peers [2]. This early candidate can then secure the majority vote and broadcast heartbeats to establish its leadership before a split vote can manifest [2].

References

- [1] Diego Ongaro and John Ousterhout. In search of an understandable consensus algorithm. In *2014 USENIX Annual Technical Conference (USENIX ATC 14)*, pages 305–319, Philadelphia, PA, June 2014. USENIX Association.
- [2] Andrew S. Tanenbaum and Maarten van Steen. *Distributed Systems: Principles and Paradigms*. Pearson Prentice Hall, 2nd edition, 2007. 686 pages.